# CPSC 322: Introduction to Artificial Intelligence

#### Uninformed Search

Textbook reference: [3.5]

Instructor: Varada Kolhatkar University of British Columbia

Credit: These slides are adapted from the slides of the previous offerings of the course. Thanks to all instructors for creating and improving the teaching material and making it available!

#### Announcements

- Please take the midterm timing survey on Canvas.
- Assignment 1 has been released.
  - We have also released the .docx version so that it's easier for editing.
- We expect to return assignment0 grades by the end of this week.

#### Lecture outline

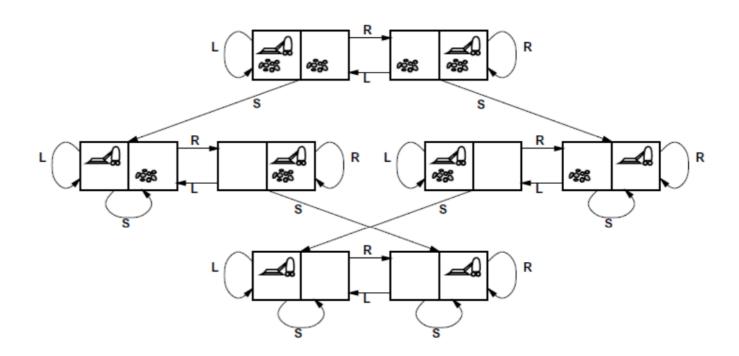
Recap from last lecture (~10 mins)



- Criteria to compare search strategies (~10 mins)
- Depth-first search (~15 mins)
- Breadth-first search (~15 mins)
- Break (~5 mins)
- Class activity (~15 mins)
- Summary and wrap-up (~5 mins)

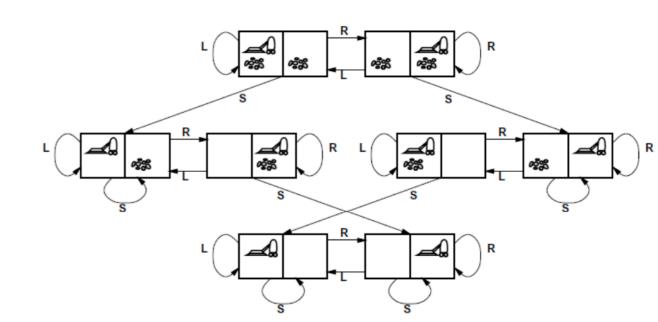
## State space graphs

- We formulated problems using
  - States
  - Initial state
  - Actions (operators)
  - Transition model
  - Goal state



## State space graphs

- State space graph is a mathematical representation of a search problem
  - Nodes are abstracted world configurations
  - Arcs represent action results
  - A goal is a set of goal nodes.
- In a search graph, each state occurs only once!



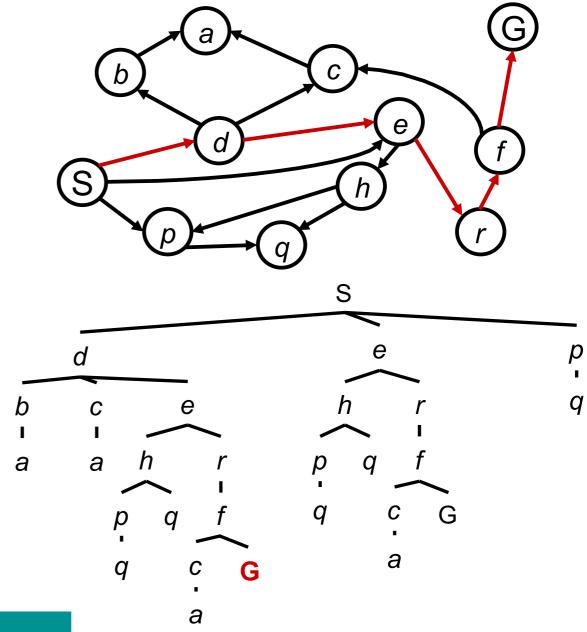
We can rarely build this graph in memory because its too big. But it's a useful idea.

#### Search trees

- Once the problem is formulated, we need to solve it.
- A solution is an action sequence.
- We find solutions using search algorithms, which consider various possible action sequences with the help of search trees.
- The possible sequences starting at the initial state form a search tree with the initial state at the root; the branches are actions, and the nodes correspond to states in the state space of the problem.

#### Search trees

- Search tree is a "what if" tree of plans and their outcomes.
- Nodes show states, but correspond to plans that achieve those states; each node in the search tree is an entire path in the state space graph.



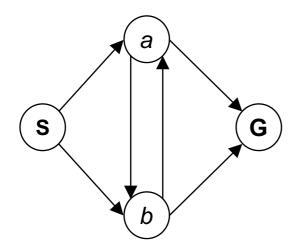
For most problems, we can never actually build the whole tree.

Credit: Berkeley Al course material

#### State space graphs vs. search trees

Consider this state graph

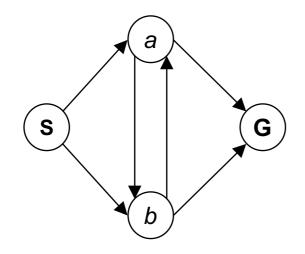
How big is the search tree?



#### State space graphs vs. search trees

Consider this state graph

How big is the search tree?





Lots of repeated structure in the search tree.

#### Activity: Find three bugs in the generic search algorithm below

```
Inputs: a graph,
          a start node no
          a boolean procedure goal(n) that tests if n is a goal node
frontier:= [ < g >: g \text{ is a start node}];
While frontier is not empty:
    select and remove path \langle n_0, n_1, \dots, n_k \rangle from frontier;
    If goal(n_k)
          return < n_0, n_1, ..., n_k >;
    Find a neighbour n of n_k
          add \langle n \rangle to frontier;
return NULL
```

#### Activity: Find three bugs in the generic search algorithm below

```
Inputs: a graph,
          a start node no
          a boolean procedure goal(n) that tests if n is a goal node
frontier:= [ < g >: g \text{ is a start node}];
While frontier is not empty:
    select and remove path \langle n_0, n_1, \dots, n_k \rangle from frontier;
    If goal(n_k)
          return < n_0, n_1, ..., n_k >;
    Find a neighbour n of n_k
          add \langle n \rangle to frontier;
return NULL
```

## Generic search algorithm

```
Inputs: a graph,
          a start node no
          a boolean procedure goal(n) that tests if n is a goal node
frontier:= [ \langle n_0 \rangle: n_0 is a start node];
While frontier is not empty:
    select and remove path \langle n_0, n_1, \dots, n_k \rangle from frontier;
    If goal(n_k)
          return < n_0, n_1, ..., n_k >;
    For every neighbour n of n_k
          add \langle n_0, n_1, \dots, n_k, n \rangle to frontier;
return NULL
```

#### Learning outcomes

- Define basic properties of search algorithms:
  - Completeness, optimality, time and space complexity
- Select the most appropriate search algorithms for specific problems.
- BFS vs. DFS vs. IDS
- LCFS vs. BestFS
- A\* vs. B&B vs. IDA\* vs. MBA\*

#### Lecture outline

- Recap from last lecture (~10 mins)
- Criteria to compare search strategies (~10 mins)



- Depth-first search (~15 mins)
- Breadth-first search (~15 mins)
- Break (~5 mins)
- Class activity (~15 mins)
- Summary and wrap-up (~5 mins)

#### Measuring problem-solving performance

We can evaluate search algorithm's performance in four ways:

- Completeness
- Optimality
- Time complexity
- Space complexity

#### Completeness and optimality

A search algorithm is **complete** if, whenever at least one solution exists, the algorithm is guaranteed to find a solution within a finite amount of time.

A search algorithm is **optimal** if, when it returns a solution, it is the best solution (i.e., there is no better solution).

Optimality can be in terms of path costs, for example.

## Time and space complexity

The **time complexity** of a search algorithm is an expression for the worst-case amount of time it will take to run expressed in terms of the **maximum path length** m and the **maximum branching factor** b.

The **space complexity** of a search algorithm is an expression for the worst-case amount of memory that the algorithm will use (number of paths). Also expressed in terms of the **maximum path length** m and the **maximum branching factor** b.

Why not in terms of |V| and |E|?

#### Lecture outline

- Recap from last lecture (~10 mins)
- Criteria to compare search strategies (~10 mins)
- Depth-first search (~15 mins)
- Breadth-first search (~15 mins)
- Break (~5 mins)
- Class activity (~15 mins)
- Summary and wrap-up (~5 mins)

In DFS the frontier is a last-in-first-out stack.

**Inputs:** a graph, a start node n<sub>o</sub>

a boolean procedure *goal(n)* that tests if

n is a goal node

frontier:= [  $\langle n_0 \rangle$ :  $n_0$  is a start node];

While frontier is not empty:

**select** and **remove** path  $\langle n_0, n_1, \dots, n_k \rangle$  from frontier;

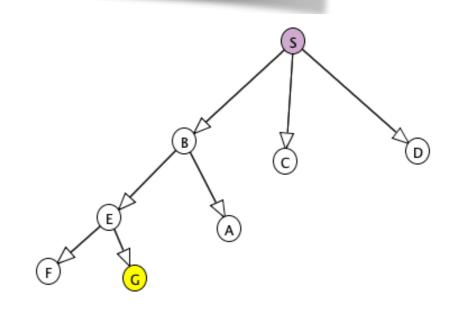
If  $goal(n_k)$ 

**return** <  $n_0, n_1, ..., n_k >$ ;

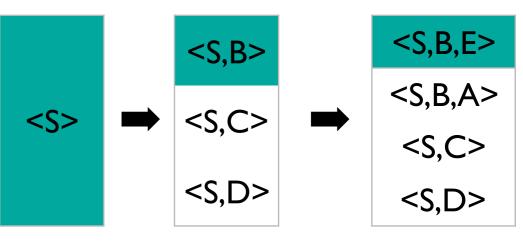
For every neighbour n of  $n_k$ 

add  $\langle n_0, n_1, \dots, n_k, n \rangle$  to frontier;

return NULL

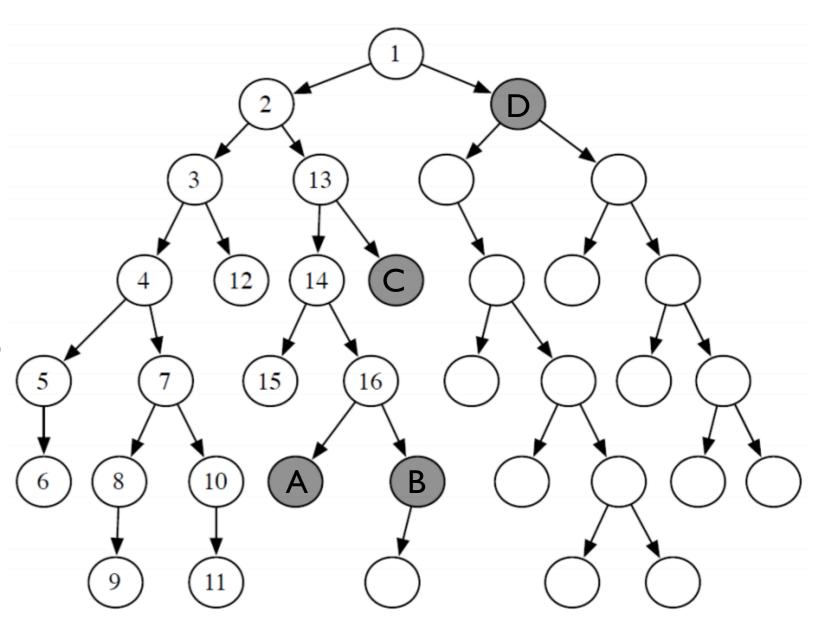


Top of the stack



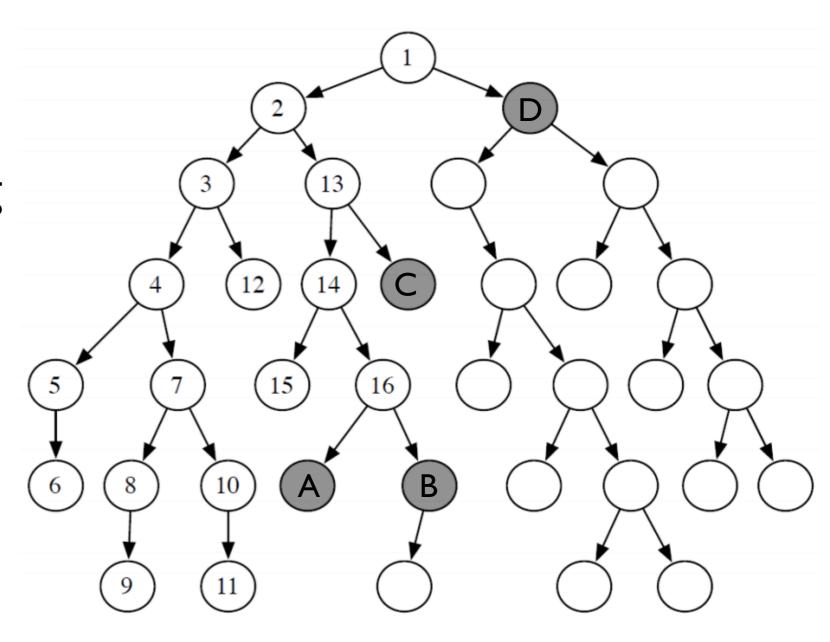
 Expands the deepest node in the current frontier of the search tree

What's frontier here?

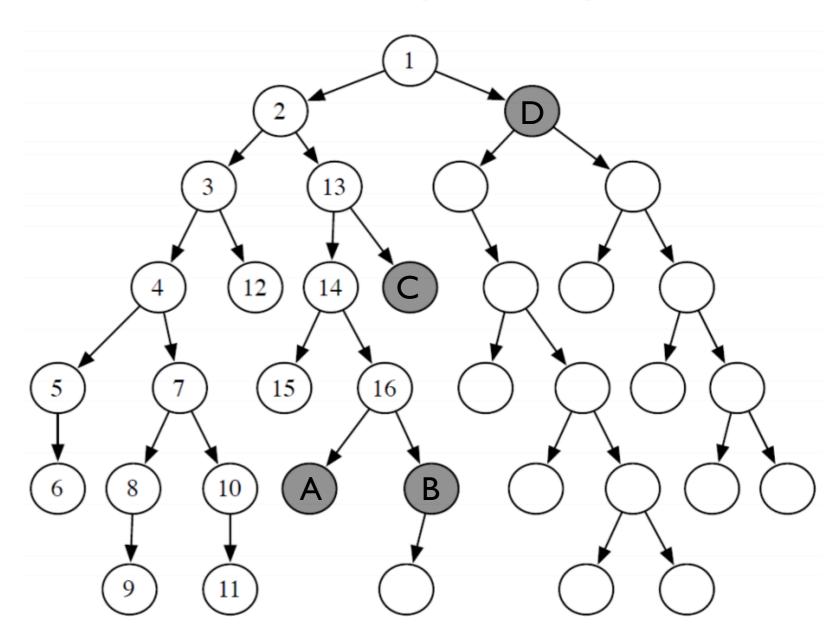


• Frontier?

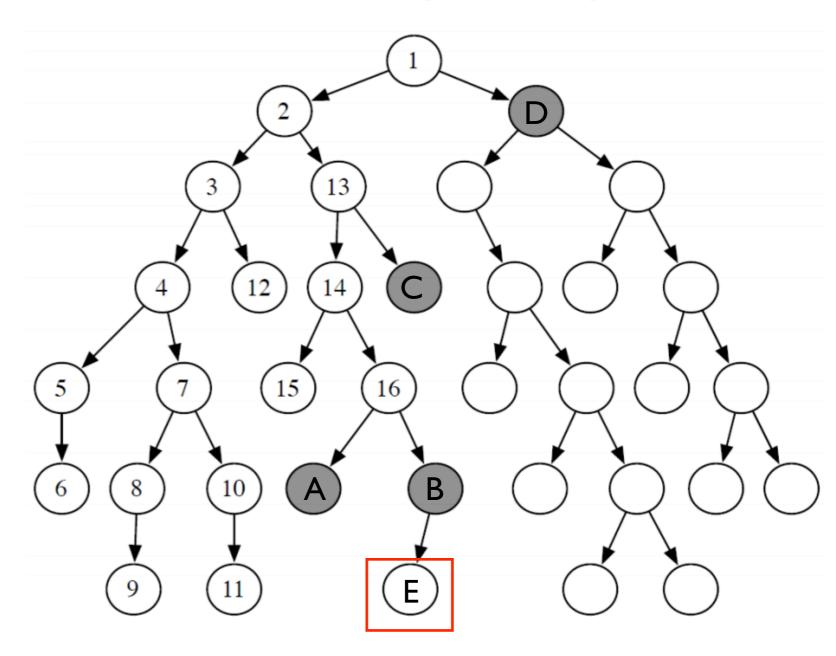
 Paths representing the shaded nodes



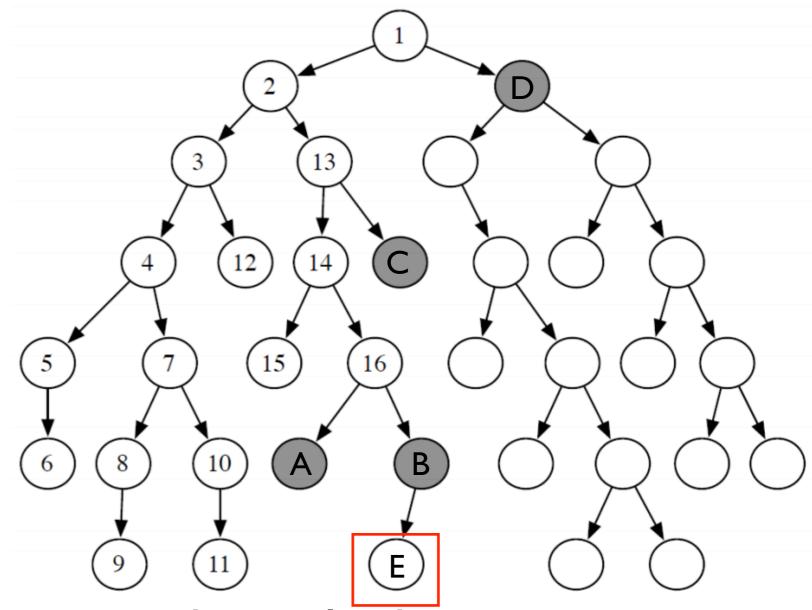
- Frontier?
  - Shaded nodes
- Which node will be expanded next?
- expand = "remove path ending at node from frontier & put its successors on"



- Say, node in red box is a goal
- How many more nodes will be expanded?



- Say, node in red box is a goal
- How many more nodes will be expanded?



3 because you only return once the goal is being expanded and not when a goal is put onto the frontier

## Analysis of DFS

## Completeness of DFS isclicker.

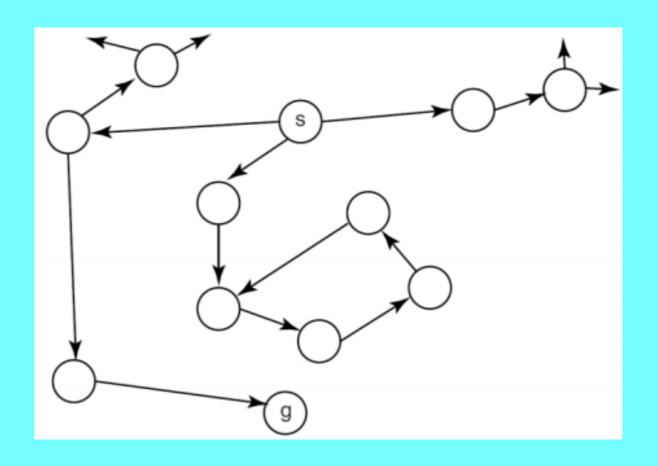


A search algorithm is complete if, whenever at least one solution exists, the algorithm is guaranteed to find a solution within a finite amount of time.

Is DFS complete?

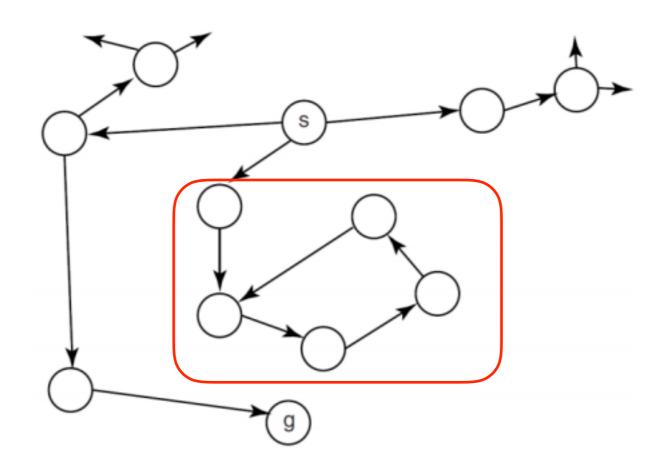
A. Yes

B. No 🔽



#### Completeness of DFS

If there are cycles in the graph, DFS might get "stuck" in one of them.



## Optimality of DFS isclicker.

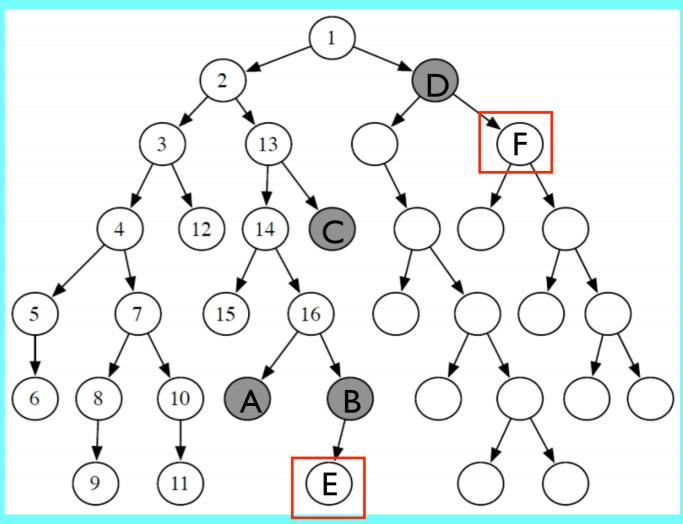


A search algorithm is optimal if, when it returns a solution, it is the best solution (i.e. there is no better solution).

Is DFS optimal?

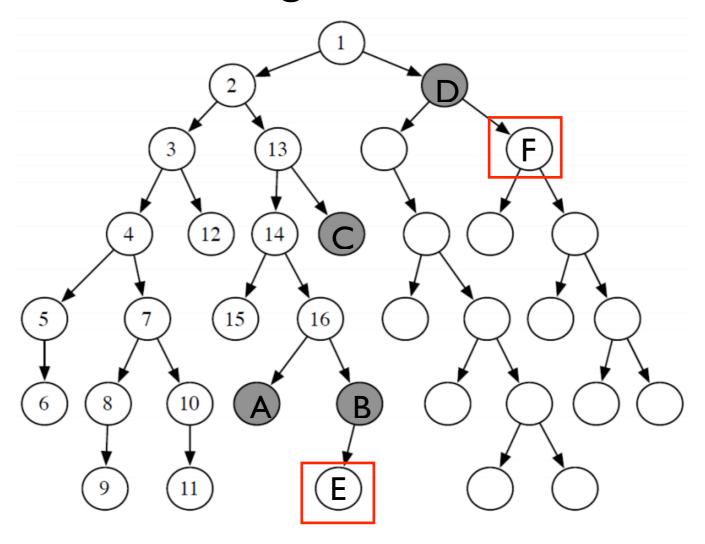
A. Yes

B. No V



#### Optimality of DFS

It can "stumble" onto longer solution paths before it gets to shorter one.



## Time complexity of DFS

iclicker.

The **time complexity** of a search algorithm is an expression for the worst-case amount of time it will take to run expressed in terms of the maximum path length m and the maximum branching factor b

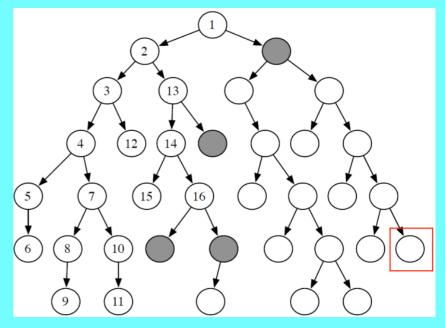
What's the time complexity of DFS in terms of m and b?

A.  $O(b^m)$ 

C. O(mb)

B.  $O(m^b)$ 

D. O(b+m)



Single goal node

## Time complexity of DFS

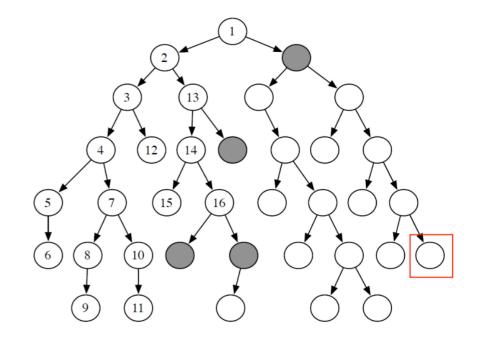
What's the time complexity of DFS in terms of m and b?

A. 
$$O(b^m)$$

 $\mathsf{C}.\ O(mb)$ 

B. 
$$O(m^b)$$

D. O(b+m)



Single goal node

In the worst case, it must examine every node in the search tree.

#### Space complexity of DFS

iclicker.

The **space complexity** of a search algorithm is an expression for the worst-case amount of memory that the algorithm will use (number of paths), expressed in terms of the maximum path length m and the maximum branching factor b.

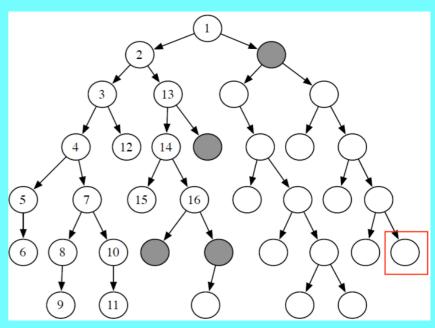
What's the space complexity of DFS in terms of m and b?

A.  $O(b^m)$ 

C. O(mb)

B.  $O(m^b)$ 

D. O(b+m)



Single goal node

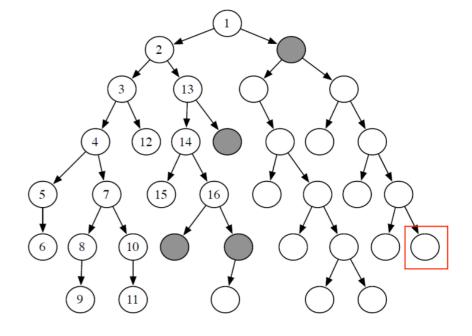
## Space complexity of DFS

What's the space complexity of DFS in terms of m and b?

A. 
$$O(b^m)$$

B. 
$$O(m^b)$$

D. 
$$O(b+m)$$



Single goal node

The longest possible path is m, and for every node in that path must maintain a fringe of size b.

## Summary of DFS analysis

- Is DFS complete? No
  - May not halt on graphs with cycles. However, DFS is complete for finite acyclic graphs.
- Is DFS optimal? No
  - It may stumble on a suboptimal solution first
- What is the time and space complexity, if the maximum path length is *m* and the maximum branching factor is *b*?
  - Time complexity is  $b^m$ : may need to examine every node in the tree.
  - Space complexity is bm: the longest possible path is m, and for every node in that path we must maintain a "fringe" of size b.

## Why bother understanding DFS?

- It is simple enough to allow you to learn the basic aspects of searching (when compared with breadthfirst search)
- It is the basis for a number of more sophisticated and/ or useful search algorithms

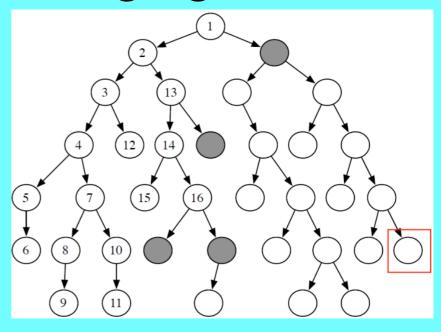
#### **DFS**



#### DFS is appropriate when

- A. There are cycles
- B. There are shallow solutions
- C. You care about optimality
- D. Space is restricted (complex space representation, e.g., in robotics)
- E. You are studying for CPSC 322

#### Single goal node



## DFS: When is it appropriate

#### **Appropriate**

- Space is restricted (complex state representation e.g., robotics assembly)
- There are many solutions, perhaps with long path lengths, particularly for the case in which all paths lead to a solution

#### Inappropriate

- When there are cycles
- When there are shallow solutions
- If you care about optimality

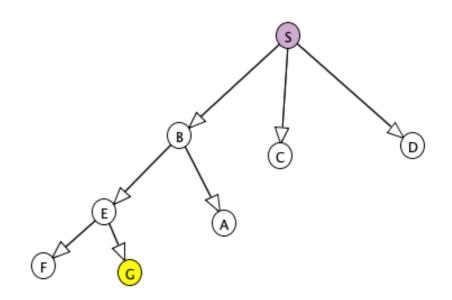
## Lecture outline

- Recap from last lecture (~10 mins)
- Criteria to compare search strategies (~10 mins)
- Depth-first search (~15 mins)
- Breadth-first search (~15 mins)
- Break (~5 mins)
- Class activity (~15 mins)
- Summary and wrap-up (~5 mins)

#### BFS as an instantiation of generic search algorithm

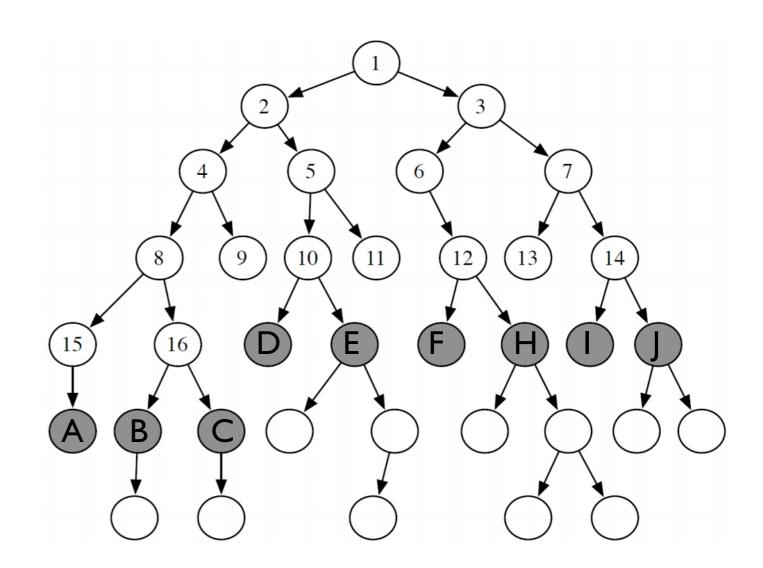
In BFS the **frontier** is a **first-in-first-out queue**.

```
Inputs: a graph,
          a start node n<sub>o</sub>
          a boolean procedure goal(n) that tests if
   n is a goal node
frontier:= [ \langle n_0 \rangle: n_0 is a start node];
While frontier is not empty:
    select and remove path \langle n_0, n_1, \dots, n_k \rangle
   from frontier;
    If goal(n_k)
          return < n_0, n_1, \dots, n_k > ;
     For every neighbour n of n_k
           add \langle n_0, n_1, \dots, n_k, n \rangle to frontier;
return NULL
```



# Breadth-first search (BFS)

- The shallowest unexpanded node is chosen for expansion.
- All nodes are expanded at a given depth in the search tree before any nodes at the next level are expanded.



# Analysis of BFS

# Completeness of BFS iclicker.



A search algorithm is complete if, whenever at least one solution exists, the algorithm is guaranteed to find a solution within a finite amount of time.

Is BFS complete?

A. Yes

B. No

# Optimality of BFS isclicker.



A search algorithm is optimal if, when it returns a solution, it is the best solution (i.e. there is no better solution)

Is BFS optimal?

A. Yes

B. No

# Time complexity of BFS

iclicker.

The **time complexity** of a search algorithm is an expression for the worst-case amount of time it will take to run expressed in terms of the maximum path length m and the maximum branching factor b.

What's the time complexity of BFS in terms of m and b?

A. 
$$O(b^m)$$

$$\mathsf{C}.\ \mathit{O}(mb)$$

B. 
$$O(m^b)$$

D. 
$$O(b+m)$$

# Space complexity of BFS

iclicker.

The **space complexity** of a search algorithm is an expression for the worst-case amount of memory that the algorithm will use (number of paths), expressed in terms of the maximum path length m and the maximum branching factor b.

What's the space complexity of BFS in terms of m and b?

A. 
$$O(b^m)$$

$$\mathsf{C}.\ O(mb)$$

B. 
$$O(m^b)$$

D. 
$$O(b+m)$$

# Summary of BFS analysis

- Is BFS complete? Yes
  - Does not get stuck in cycles
- Is BFS optimal? Yes
  - Guaranteed to find the path that involves fewest arcs
- What is the time and space complexity, if the maximum path length is *m* and the maximum branching factor is *b*?
  - Time complexity is  $b^m$ : may need to examine every node in the tree.
  - Space complexity is  $b^m$ : frontier contains all paths of the relevant length (which is  $\leq$  the shortest path length to a goal node)

## Time and memory requirements of BFS

The numbers in the table assume branching factor b=10; a million nodes/second and 1000 byte/node

Depth	Nodes	Time	Memory
2	$10^{2}$	.11 ms	107 KB
4	$10^{4}$	II ms	10.6 MB
6	$10^{6}$	I.I sec	I GB
8	$10^{8}$	2 mins	103 GB
10	$10^{10}$	3 hours	IOTB
12	$10^{12}$	13 days	I Petabyte
14	$10^{14}$	3.5 years	99 petabyte
16	$10^{16}$	350 years	10 exabyte

Source: Russell and Norvig Figure 3.13

## BFS: When is it appropriate?

#### **Appropriate**

- Space is not a problem
- It's necessary to find the solution with the fewest arcs
- Although all solutions may not be shallow, at least some are

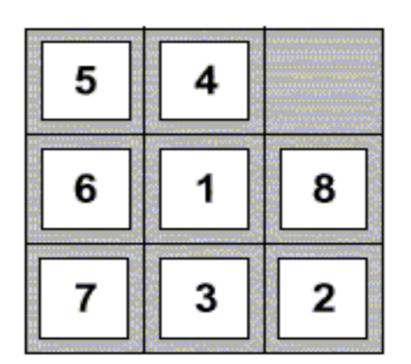
#### Inappropriate

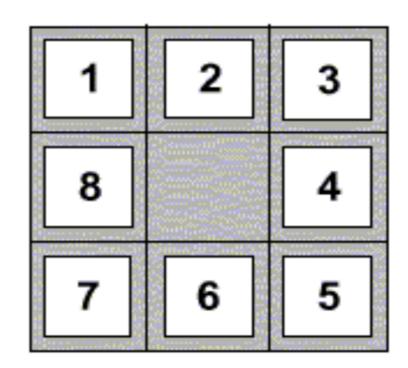
- Space is limited
- All solutions tend to be located deep in the tree
- The branching factor is very large

# When to use BFS vs. DFS A. BFS B. DFS

- I. The search graph has cycles or is infinite.
- 2. We need a shortest path to a solution.
- 3. There are only solutions at great depth.
- 4. There are some solutions at shallow depth.
- 5. Memory is limited.

# Real example: Solving Eight Puzzle





Which search method would you use if you want to find the shortest solution?

A. BFS B. DFS

## What have we done so far?

- Learned how to formulate a problem as a search problem with an initial state, a goal state, a set of actions, and a transition model.
- Studied the generic search algorithm. Studied two instantiations of the generic search algorithm, BFS and DFS, which can find solutions (plans) given a search problem.
- Compared the algorithms using four different measures: completeness, optimality, time and space complexity.

# Summary table

	complete?	optimal?	time O()	space O()
DFS	FALSE	FALSE	$b^m$	bm
BFS	TRUE	True*	$b^m$	$b^m$

<sup>\*</sup>Assuming arcs all have the same cost. (We'll get to this later.)

# Class activity (~15 mins)

### Homework

- To test your understanding of today's class
  - Work on Practice Exercise 3.B on <a href="http://www.aispace.org/exercises.shtml">http://www.aispace.org/exercises.shtml</a>

# Coming up

- Read the following before the next class.
  - For the next Iterative deepening [3.5.3]
  - Search with costs [3.5.4]
  - Heuristic Search [3.6]